

Rifat Erdem Sahin — Kong AI Gateway CV

RIFAT ERDEM SAHIN

API Platform & AI Gateway Architect — Kong | Multi-LLM Orchestration

 **Location:** London, United Kingdom  **Citizenship:** British  **Email:** contact@rifaterdemsahin.com  **Phone:** +44 7848 024173  **LinkedIn:** [linkedin.com/in/rifaterdemsahin](https://www.linkedin.com/in/rifaterdemsahin)  **GitHub:** github.com/rifaterdemsahin  **Portfolio:** <https://rifaterdemsahin.com>  **Schedule a Call:** <https://calendly.com/rifaterdem/schedule>

PROFESSIONAL SUMMARY

API Platform and AI Gateway Architect with **12+ years of experience** designing scalable, secure, and cost-optimized integration platforms. Specialist in **Kong Gateway and Kong AI Gateway**, building production-grade LLM orchestration layers that route traffic intelligently across multiple model providers. Designed and delivered the “**Smart-Cost**” **Multi-LLM Orchestrator** — an intelligent gateway that automatically routes prompts to the most cost-effective model based on request complexity, **cutting LLM API spend by 40–60%** with zero changes to the frontend application. Deep expertise in semantic routing, API governance, observability, and OpenAI-compatible API standardization.

CORE COMPETENCIES

 **Kong Gateway & Kong AI Gateway** - Kong Gateway (OSS & Enterprise) plugin architecture, routing, and policy design - Kong AI Gateway: `ai-proxy`, `ai-proxy-advanced`, and `ai-semantic-router` plugins - Semantic routing based on prompt intent, complexity, and embeddings - Multi-LLM load balancing, fallback chains, and provider failover - Declarative configuration with `deck`, GitOps, and Kong Konnect

 **Multi-LLM Orchestration** - OpenAI-compatible API standardization across heterogeneous providers - Cost-aware model selection (gpt-4o-mini, gemini-flash, claude-3-5-sonnet) - Token budgeting, rate limiting, and per-tenant quota enforcement - Prompt caching, response streaming, and latency optimization

 **API Management & Integration** - API lifecycle management, versioning, and developer portals - Authentication (OAuth2, JWT, API keys), mTLS, and zero-trust gateways - Plugin development in Lua and external plugins (Go/Python/JS) - Service mesh integration with Kong Mesh and Kuma

 **Observability & FinOps** - LLM cost telemetry, per-model dashboards, and spend attribution - Prometheus, Grafana, OpenTelemetry tracing across the gateway - SLA monitoring, request analytics, and anomaly detection

☁️ **Cloud & DevOps** - Kubernetes-native deployment of Kong Ingress Controller - Terraform, Helm, CI/CD pipelines on Azure, AWS, and GCP - Containerization with Docker, GitOps with ArgoCD

KEY ACCOMPLISHMENTS

🏆 2025 | Flagship Project | The “Smart-Cost” Multi-LLM Orchestrator

An intelligent Kong AI Gateway that slashes LLM bills by dynamically matching model strength to task complexity — completely invisible to the frontend.

The Concept Built an intelligent gateway that automatically routes user prompts to the most cost-effective LLM based on the difficulty of each request, demonstrating how a business can dramatically reduce LLM API bills without touching the client application.

How It Works - Delivered a simple user chat interface that sends **every** request to Kong using a single, standardized **OpenAI-compatible API format** — the frontend never knows which model answers. - Kong inspects each prompt using **semantic routing** and classifies it by complexity. - **Lightweight requests** (e.g., “*Summarize this paragraph*”) are routed to a fast, inexpensive model such as **gpt-4o-mini** or **gemini-flash**. - **Complex requests** (e.g., “*Debug this Python script*”) are routed to a heavy-duty model such as **claude-3-5-sonnet**.

Kong’s Role - Configured the **Kong AI Gateway** semantic router to inspect prompt content and embeddings, then select the optimal upstream provider in real time. - Normalized all upstream providers behind one OpenAI-compatible contract so the chat client uses a single endpoint and schema. - Added fallback chains, rate limiting, and per-model cost telemetry for full spend visibility.

Why It’s Valuable - Demonstrated that businesses can **slash LLM API bills by dynamically matching the model’s strength to the task complexity**, completely invisible to the frontend application. - Achieved **40–60% cost reduction** on representative production traffic while preserving answer quality on complex tasks. - Established a reusable reference architecture for cost-aware AI gateways adopted across multiple teams.

Tech Stack: Kong AI Gateway · ai-proxy-advanced · ai-semantic-router · OpenAI-compatible API · gpt-4o-mini · gemini-flash · claude-3-5-sonnet · Kubernetes · Prometheus · Grafana

🏆 2024 | Enterprise API Platform Modernization

Kong Gateway Migration & Governance at Scale - Migrated 200+ legacy APIs onto Kong Gateway with declarative **deck** configuration and GitOps workflows. - Implemented centralized authentication, rate limiting, and observability, reducing onboarding time for new services by **70%**. - Standardized plugin policies across environments, cutting production incidents related to misconfiguration by **45%**.

🏆 2023 | AI Application Delivery

Production RAG & LLM Systems Behind a Unified Gateway - Architected retrieval-augmented generation services fronted by an API gateway with prompt-injection protection and PII redaction. -

Introduced token budgeting and caching, reducing redundant model calls by **35%**.

TECHNICAL SKILLS

API Gateways: Kong Gateway (OSS/Enterprise), Kong AI Gateway, Kong Konnect, Kong Ingress Controller, Kong Mesh **AI/LLM:** OpenAI API, Anthropic Claude, Google Gemini, semantic routing, embeddings, RAG, prompt engineering **Languages:** Lua (Kong plugins), Python, Go, JavaScript/TypeScript, Bash **Infrastructure:** Kubernetes, Docker, Terraform, Helm, ArgoCD **Observability:** Prometheus, Grafana, OpenTelemetry, ELK **Cloud:** Azure, AWS, GCP **Tooling:** deck, GitOps, CI/CD (GitHub Actions, Azure DevOps)

CERTIFICATIONS & EDUCATION

- Kong Certified — AI Gateway & API Management practices
 - Cloud architecture certifications across Azure / AWS
 - BSc, Computer Engineering
-

WHY KONG

Kong sits at the perfect intersection of **API management** and the **AI revolution**. The Kong AI Gateway turns the hardest problem in enterprise AI adoption — controlling cost, governance, and provider sprawl — into a configuration concern. I want to help businesses prove that intelligent, semantic routing can make LLM spend predictable and frontend-invisible, exactly as demonstrated by the Smart-Cost Multi-LLM Orchestrator.

References and live demo of the Smart-Cost Multi-LLM Orchestrator available on request.